



Azure Virtual Machines (VMs)

◆ What It Is

Azure Virtual Machines (VMs) provide Infrastructure as a Service (IaaS), allowing you to run Windows or Linux virtual machines in the cloud. VMs are highly customizable and behave much like physical servers, giving full control over the operating system, installed software, and network configuration. They are ideal for lift-and-shift scenarios, custom software stacks, legacy workloads, and full administrative access.

◆ How It Works / How to Implement on Azure

To deploy a VM in Azure, you typically go through the following steps:

- **Choose an image:** Select a base OS image from the Azure Marketplace (e.g., Windows Server, Ubuntu).
 - **Select a size:** VM sizes (SKUs) define CPU, RAM, and IOPS capabilities.
 - **Create or select a resource group:** Logical container for the VM and related resources.
 - **Configure networking:** Assign the VM to a Virtual Network (VNet), with a subnet, public IP (optional), and Network Security Group (NSG).
 - **Set storage:** Choose OS disk and optional data disks (HDD/SSD, standard/premium).
 - **Authentication:** SSH key or username/password for login.
 - **Deployment tools:** You can create VMs using the Azure Portal, Azure CLI, PowerShell, ARM templates, or Bicep.
-

◆ Pros and Cons

Pros:

- Full control over OS and installed software.
- Flexible in terms of configurations and scaling.
- Supports both Linux and Windows workloads.
- Ideal for custom, legacy, or resource-intensive apps.

Cons:

- More maintenance: OS patching, antivirus, monitoring.
- Higher operational overhead compared to PaaS.

- Slower scaling compared to containerized or serverless options.
 - Pricing can grow quickly if not monitored.
-

◆ Key Considerations

- **Availability:** Use *Availability Sets* or *Availability Zones* for high availability.
 - **Disaster Recovery:** Consider Azure Site Recovery or VM backups.
 - **Security:** Harden the OS, use NSGs, Azure Bastion, and consider Just-In-Time access.
 - **Performance:** Choose SSD-based disks for better IOPS; use Accelerated Networking for high-throughput workloads.
 - **Dev/Test vs. Production:** Use *Dev/Test Pricing* for eligible subscriptions, or Spot VMs for ephemeral workloads.
-

◆ Pricing, Consumption & Scaling Tips

- **Pricing model:** Pay per second while the VM is running; charges include compute, OS disk, and networking.
- **Scaling:** Manual or automated scaling with **VM Scale Sets**. Auto-scale based on metrics like CPU or custom rules.
- **Cost saving tips:**
 - Use *Azure Reservations* for long-term workloads (1-year or 3-year reserved instances).
 - Use *Spot VMs* for interruptible workloads at a significant discount.
 - Deallocate VMs when not in use to stop charges (but keep storage).
 - Monitor with Azure Cost Management and set budgets or alerts.



Azure App Services

◆ What It Is

Azure App Services is a **Platform as a Service (PaaS)** offering that enables you to build, host, and scale web apps, RESTful APIs, and mobile backends quickly without managing infrastructure. It supports multiple languages and frameworks including .NET, Java, Python, Node.js, PHP, and Ruby. It's designed to let developers focus on app logic while Azure handles OS patches, scaling, and load balancing.

◆ How It Works / How to Implement on Azure

To deploy an app using App Services:

- **Create an App Service Plan:** Defines the region, compute resources (CPU/RAM), and pricing tier.
- **Deploy the Web App:** Choose a runtime stack and publish your app (via GitHub, DevOps, ZIP deployment, FTP, or local Git).
- **Configure settings:** Set environment variables, scaling rules, custom domains, SSL, authentication, and more.
- **Use deployment slots:** For staging/production swap or A/B testing.

Management is done through the **Azure Portal**, **CLI**, **ARM templates**, **Bicep**, or tools like **Visual Studio** and **VS Code** with Azure extensions.

◆ Pros and Cons

Pros:

- Zero infrastructure management.
- Built-in CI/CD and deployment slots.
- Auto-scaling and load balancing out of the box.
- Integrates easily with Azure services (e.g., Azure SQL, Key Vault).
- Supports custom domains and SSL.

Cons:

- Less control compared to VMs (no access to underlying OS).
- Cold starts and scaling delays in lower pricing tiers.
- Limited to supported runtimes and configurations.

- Higher tiers required for VNet integration or more advanced features.
-

◆ Key Considerations

- **App Service Plan choice** impacts performance and price — choose wisely based on CPU/RAM needs.
 - **Scaling:** Horizontal (instances) and vertical (pricing tier) scaling available.
 - **Security:** Use Azure App Service Authentication (Easy Auth), integrate with Azure AD, and configure private endpoints for VNet integration.
 - **Dev/Test vs. Production:** Free and Shared tiers are good for learning/testing but have limits. For production, consider Standard or Premium tiers.
 - **CI/CD:** Seamless integration with GitHub Actions, Azure DevOps, Bitbucket, etc.
-

◆ Pricing, Consumption & Scaling Tips

- **Pricing model:** Based on the App Service Plan — you pay for the compute instance, not the individual app.
- **Free/Shared tiers:** Ideal for demos or learning, with limited compute resources.
- **Basic/Standard/Premium:** More compute, features, and scaling capacity.
- **Auto-scale:** Configure rules to scale out/in based on CPU usage, memory, or HTTP queue length.
- **Deployment Slots:** Included in Standard and Premium tiers — can save money on test environments.
- **Reserved instances:** Save on cost with long-term commitments (App Service Environment only).

Azure Kubernetes Service (AKS)

◆ What It Is

Azure Kubernetes Service (AKS) is a **managed Kubernetes** offering that lets you deploy, manage, and scale containerized applications using Kubernetes, without the overhead of managing the control plane. It abstracts much of the Kubernetes complexity, giving you powerful orchestration features for microservices, batch jobs, and stateless/stateful workloads.

◆ How It Works / How to Implement on Azure

To deploy an AKS cluster:

- **Create an AKS cluster:** Define node size/count, Kubernetes version, and networking options.
- **Choose node pools:** Can be standard VMs or spot VMs; support scaling and GPU-enabled workloads.
- **Configure networking:** Choose between Kubenet or Azure CNI for network plugins.
- **Deploy your containers:** Use `kubectl`, Helm charts, or GitOps (via Azure Arc or Flux).
- **Integrate with container registries:** Typically Azure Container Registry (ACR) or Docker Hub.

AKS handles:

- Control plane (API server, etcd, scheduler) — maintained by Azure.
 - You manage the **worker nodes** and deployed workloads.
-

◆ Pros and Cons

Pros:

- Fully managed Kubernetes control plane.
- Native integration with Azure Monitor, ACR, Key Vault, and more.
- Auto-scaling, self-healing, rolling upgrades.
- Ideal for microservices and large-scale containerized architectures.

Cons:

- Steeper learning curve than App Services or Functions.
- Still requires cluster/node management and understanding Kubernetes concepts.
- Costs can escalate quickly with complex workloads and unused nodes.

- Cold starts and provisioning time for new nodes can affect performance.
-

◆ Key Considerations

- **Scaling:** Use Cluster Autoscaler (adds/removes nodes) and Horizontal Pod Autoscaler (adjusts pod count).
 - **Security:** Apply RBAC, use private clusters, integrate Azure AD, and secure container images.
 - **Networking:** CNI gives more flexibility but uses more IPs. Consider network policies for traffic control.
 - **Node pools:** Separate pools for dev/test/prod or GPU workloads.
 - **Observability:** Use Azure Monitor for containers, Prometheus/Grafana, or third-party tools.
 - **CI/CD:** Integrate with Azure DevOps or GitHub Actions for automated deployments.
-

◆ Pricing, Consumption & Scaling Tips

- **Pricing model:** Control plane is free; you pay for the VM instances in node pools and associated resources (disks, load balancers, networking).
- **Optimize costs:**
 - Use **spot instances** for non-critical or short-lived workloads.
 - Configure autoscaling to shut down idle nodes.
 - Use smaller node sizes or mix node types to balance cost and performance.
 - Run multiple apps per node (where feasible) to improve utilization.
- **Bursting:** Use Virtual Node integration (with Azure Container Instances) for rapid scaling needs.

Azure Functions

◆ What It Is

Azure Functions is a **serverless compute service** that lets you run small pieces of code — “functions” — in the cloud without provisioning or managing infrastructure. It’s event-driven, meaning your code runs in response to triggers (HTTP requests, queue messages, timers, etc.). Ideal for lightweight logic, background tasks, and automation.

◆ How It Works / How to Implement on Azure

To create a function:

- **Choose a hosting plan:** Consumption, Premium, or Dedicated (App Service Plan).
- **Pick a runtime:** Supports C#, JavaScript, Python, PowerShell, Java, and others.
- **Define a trigger:** Common types include HTTP, Timer, Blob Storage, Azure Queue, Event Grid, and Service Bus.
- **Write your function code:** In the Azure Portal, VS Code, or locally using the Azure Functions Core Tools.
- **Deploy:** Via GitHub Actions, Azure DevOps, ZIP deploy, or Azure CLI.

Functions scale automatically, run statelessly, and can integrate with other Azure services using **bindings** — a powerful way to connect inputs and outputs without writing boilerplate code.

◆ Pros and Cons

Pros:

- Extremely fast to get started.
- No infrastructure management.
- Scales automatically — from zero to thousands of concurrent executions.
- Cost-effective for intermittent workloads.
- Rich integration with Azure and third-party services via bindings.

Cons:

- Cold starts in the Consumption Plan may delay execution (esp. in low-traffic apps).
- Stateless by design — needs external storage (e.g., Azure Storage or Cosmos DB) for state.
- Timeouts (default 5 min in Consumption Plan, extendable to 60 min with Premium).

- Vendor lock-in risk due to Azure-specific features.
-

◆ Key Considerations

- **Hosting plan:**
 - **Consumption Plan:** Pay-per-execution, great for low-traffic or bursty workloads.
 - **Premium Plan:** No cold starts, VNET integration, longer execution time.
 - **Dedicated (App Service Plan):** Useful if co-hosting with other web apps.
 - **State:** Keep it external; use Durable Functions if orchestration or chaining is needed.
 - **Error handling:** Implement retries and dead-letter queues for robustness.
 - **Monitoring:** Use Application Insights for real-time metrics, logs, and traces.
 - **Security:** Secure HTTP triggers with function keys, tokens, or Azure AD.
-

◆ Pricing, Consumption & Scaling Tips

- **Consumption Plan:**
 - **Billing:** Based on GB-seconds and number of executions.
 - Ideal for event-driven tasks like data processing, backend jobs, or light APIs.
 - **First 1M executions and 400,000 GB-s free per month.**
- **Premium Plan:**
 - Reserved instances (even for idle time), but eliminates cold starts and supports VNET.
- **Scaling:**
 - Automatic scaling based on event volume.
 - Use Durable Functions to handle orchestrations, retries, and stateful flows.
- **Optimize costs:**
 - Use bindings to avoid writing extra integration code (e.g., direct output to Blob Storage).
 - Limit concurrency when needed using host.json settings.



Subscriptions

An **Azure Subscription** is the unit of billing and resource organization in Azure. Every resource you create in Azure is tied to a subscription, which acts as a **container for billing, permissions, and policies**. You can think of it like a boundary that defines what resources you're charged for, and under what terms (e.g. pay-as-you-go, free trial, enterprise agreements).

Subscriptions are linked to **Azure Active Directory (AAD)** tenants for identity and access control. Within a tenant, you can have multiple subscriptions — useful for separating environments (e.g., dev, test, prod), departments, or cost centers. Role-based access control (RBAC) can be applied at the subscription level, and **Azure Policy** or **management groups** can enforce governance across multiple subscriptions.



Resource Groups

A **Resource Group (RG)** is a logical container that holds related Azure resources — like VMs, storage accounts, virtual networks, etc. It's the **deployment and management boundary** for those resources, allowing you to manage them collectively (e.g., deleting or applying tags to an entire RG).

Resource groups are **region-independent** (though the resources inside them may be tied to specific regions), and they play a key role in organizing and controlling lifecycle, access, and monitoring. You can apply RBAC and policies at the RG level, making it easier to manage access and compliance per project or team. In most cases, resources in an RG should share the same lifecycle (i.e., be created and deleted together).

Regions, Availability Zones, and Availability Sets — key concepts for understanding the geographical distribution and resilience of Azure services:

Regions

Azure Regions are physical locations around the world where Azure data centers are clustered to provide cloud services. Each region contains one or more **datacenters** offering Azure services, such as **Virtual Machines (VMs)**, **Storage**, and **Databases**. By choosing an appropriate region for your resources, you can reduce **latency** and comply with **data residency** requirements, as well as meet **regulatory** or **compliance** needs.

Azure offers **global reach**, with **over 60 regions** worldwide. The selection of a region depends on factors such as proximity to end-users, compliance, and the **availability** of specific Azure services in that location. Using regions allows you to deploy applications closer to customers for improved performance while taking advantage of the scalability and reliability of the cloud.

Availability Zones

Availability Zones are **physically separate** locations within an Azure region that provide high availability and fault tolerance. Each Availability Zone has its own **power**, **network**, and **cooling**, designed to ensure that even if one zone experiences a failure, others will continue operating normally. This makes Availability Zones ideal for **high-availability** and **disaster recovery** setups.

By deploying resources like VMs, storage, and databases across multiple Availability Zones, you can ensure your application is **resilient** to outages at the datacenter level. It's often used for critical applications that cannot afford downtime. Availability Zones are supported in **select Azure regions**, so check availability when designing your architecture.

Availability Sets

Availability Sets are a feature within an Azure region that helps **ensure application uptime** by distributing **VMs** across multiple **fault domains** and **update domains**. A **fault domain** is a group of VMs that share a common power source and network switch, while an **update domain** ensures VMs are updated in batches to avoid downtime during maintenance.

While **Availability Sets** provide high availability by ensuring your VMs are distributed across different physical resources, they don't offer the level of redundancy provided by **Availability Zones**, which are designed for more robust resilience. Availability Sets are ideal for less critical applications that require **fault tolerance** but don't need the geographical separation offered by Availability Zones.

deployment and infrastructure management tools:

ARM Templates

Azure Resource Manager (ARM) Templates are JSON files used to **define infrastructure as code**. They declaratively specify Azure resources and configurations — allowing you to **provision, update, and version** your environment consistently and repeatably.

ARM Templates support **parameterization, modularization, and conditionals**, making them powerful for deploying complex architectures (e.g., a full web app with networking, databases, and monitoring). They're idempotent — meaning you can deploy them multiple times without creating duplicate resources — and they work well with CI/CD pipelines, especially when paired with Azure DevOps or GitHub Actions.

Scale Sets

Virtual Machine Scale Sets (VMSS) let you **deploy and manage a group of identical, load-balanced VMs**. They support both Windows and Linux VMs, and automatically scale in or out based on demand (e.g., CPU usage, schedule, or custom metrics via autoscale rules).

VMSS is ideal for high-availability, stateless applications. Integrated with Azure Load Balancer and Azure Application Gateway, scale sets can help distribute traffic evenly. They also support **rolling updates, automatic OS image upgrades**, and integration with Azure Spot VMs for cost-saving scenarios. For stateful workloads, managed disks can now be attached in newer configurations.

Deployment Slots

Deployment Slots are a feature of **Azure App Service** (Web Apps) that let you **run multiple live versions of your app**, such as “staging,” “testing,” and “production.” You can swap slots with zero downtime, making this ideal for **blue-green deployments, A/B testing**, or quick rollbacks if something fails in production.

Each slot has its own configuration settings (e.g., environment variables, connection strings) and operates in isolation. The swap operation preserves warm-up state and avoids cold starts. This makes Deployment Slots a safe and fast way to ship changes — especially for customer-facing web apps or APIs.

Docker Containers and **Azure Container Registry (ACR)** — core building blocks for container-based development and deployment in Azure:

Docker Containers

Docker Containers package an application and all its dependencies into a single, lightweight unit that runs consistently across environments. Unlike traditional VMs, containers **share the host OS kernel**, making them more efficient and faster to start.

In Azure, you can run Docker containers in many ways — on **Azure Kubernetes Service (AKS)**, **Azure Container Instances (ACI)** for short-lived jobs, or in **App Services with Linux containers**. Containers support microservices architecture, simplify CI/CD workflows, and improve portability. They're ideal for rapid scaling, isolation, and testing, especially in distributed or cloud-native apps.

Azure Container Registry (ACR)

Azure Container Registry is a **private Docker-compatible registry** for storing and managing your container images. It integrates seamlessly with Azure services like AKS, ACI, App Services, and CI/CD tools like GitHub Actions or Azure DevOps.

ACR allows you to **build, tag, scan, and push images** securely within Azure, avoiding reliance on public registries like Docker Hub. Features include geo-replication across regions, content trust, image signing, and integration with Azure Active Directory for access control. You can also automate builds using **ACR Tasks**, and manage image lifecycles via retention policies or webhooks.

Durable Functions and **Azure Disk** — both key elements for specialized compute and storage in Azure:

Durable Functions

Durable Functions is an extension of Azure Functions that helps manage **long-running, stateful workflows**. While traditional Azure Functions are stateless and execute in response to triggers, Durable Functions allows you to orchestrate multiple functions, chain them together, and manage **state** between function calls.

Durable Functions use a **Durable Task Framework** to persist state automatically, so you can handle workflows like human approvals, file processing, or complex retry logic. They support patterns such as **fan-out/fan-in, async HTTP requests, timers, and sub-orchestrations**. This makes them ideal for building reliable, scalable, and fault-tolerant workflows without managing infrastructure.

Azure Disk

Azure Disk Storage provides **persistent block-level storage** for Azure Virtual Machines (VMs) and other Azure services. It offers several types of disks, including **Standard HDD, Standard SSD, and Premium SSD**, which vary in performance and cost to suit different workloads.

Azure Disks are **durable**, replicated to protect against hardware failures, and can be attached to VMs for high-performance applications. You can use **Managed Disks**, which simplifies management, as Azure handles the disk creation, scaling, and redundancy. Features like **disk snapshots, encryption** (via Azure Disk Encryption), and **disk bursting** ensure data security and scalability.

SSH and **RDP** — essential tools for remote access and management of Azure resources:



SSH

SSH (Secure Shell) is a **protocol used to securely connect to remote machines** over a network. In Azure, SSH is commonly used to access **Linux-based virtual machines (VMs)** and **containerized applications** in a secure, encrypted manner. Unlike traditional RDP, which is often used for graphical user interfaces, SSH provides command-line access, which is especially useful for managing and troubleshooting VMs or containers.

When connecting to an Azure VM, SSH allows you to **execute shell commands**, install packages, configure services, and manage the file system. Azure supports **SSH key-based authentication**, which is more secure than password-based methods. You can use Azure CLI or Azure Portal to configure SSH access and easily manage your keys or troubleshoot connection issues.



RDP

RDP (Remote Desktop Protocol) is a **Microsoft protocol** that allows you to connect to **Windows-based virtual machines (VMs)** in Azure and interact with the desktop interface as if you were physically sitting at the machine. This is particularly useful for **administrative tasks, application management**, and situations where a GUI-based interface is necessary.

Azure VMs with **Windows Server** or **Windows 10/11** are commonly accessed using RDP. You can manage network security for RDP through **Network Security Groups (NSGs)** and **Azure Bastion** for secure, no-exposure connections to VMs without needing a public IP address. RDP sessions can also be encrypted for privacy and security during the connection.

the **Cloud Service Models: IaaS, PaaS, and SaaS**. These models are foundational for understanding how Azure provides various levels of abstraction for cloud resources.

IaaS (Infrastructure as a Service)

IaaS is the most basic cloud service model, providing **virtualized computing resources** over the internet. With IaaS, you rent **virtual machines (VMs)**, **storage**, and **networking** without the need to own or manage physical hardware. Azure IaaS solutions include **Azure Virtual Machines**, **Azure Load Balancer**, and **Azure Virtual Networks**, giving users full control over their infrastructure while offloading the management of physical servers.

This model is ideal for companies that need **flexible, scalable computing power** without the overhead of maintaining physical infrastructure. IaaS is suitable for running legacy applications, experimenting with new technologies, or hosting services that require custom configurations. However, you are responsible for managing the **operating system, patching, and security**.

PaaS (Platform as a Service)

PaaS provides a higher level of abstraction than IaaS by delivering a **platform** to develop, run, and manage applications without worrying about the underlying infrastructure. Azure PaaS services include **Azure App Services**, **Azure SQL Database**, and **Azure Kubernetes Service (AKS)**, where you focus on developing your application and Azure handles the platform, scaling, and maintenance.

This model is ideal for **developers** who want to focus on writing code and building features, not managing infrastructure. PaaS services typically include **built-in developer tools**, **automated scaling**, and **integration with CI/CD pipelines**, enabling faster development and more efficient operations. While you don't have to manage servers, you do have some control over the environment, such as configuring web apps or databases.

SaaS (Software as a Service)

SaaS is the highest level of abstraction, providing fully managed applications over the internet. In SaaS, the **application, infrastructure, and platform** are all managed by the cloud provider. Examples of Azure-based SaaS solutions include **Office 365**, **Azure DevOps Services**, and **Dynamics 365**.

SaaS is ideal for businesses that need **ready-to-use applications** without the need for custom development or infrastructure management. It eliminates the need for software installation, maintenance, and updates. SaaS is perfect for productivity tools, collaboration platforms, and any business functions where the primary concern is functionality and ease of use, rather than infrastructure.